

HyperPipelining of High-Speed Interface Logic

Gregg Baeckler

02-16-2016



Define “Hyper Pipeline”?

☛ [Intel ~2000]

- A 20 stage (heavy, speed record setting) Pentium 4 pipeline

☛ [Altera internal, ~2011]

- The word assigned to the already prevalent network logic practice of adding more registers until the circuit stops going faster

☛ [Altera, now part of Intel 2015]

- Pipelining a RTL design with Stratix10 HyperFlex™ Registers

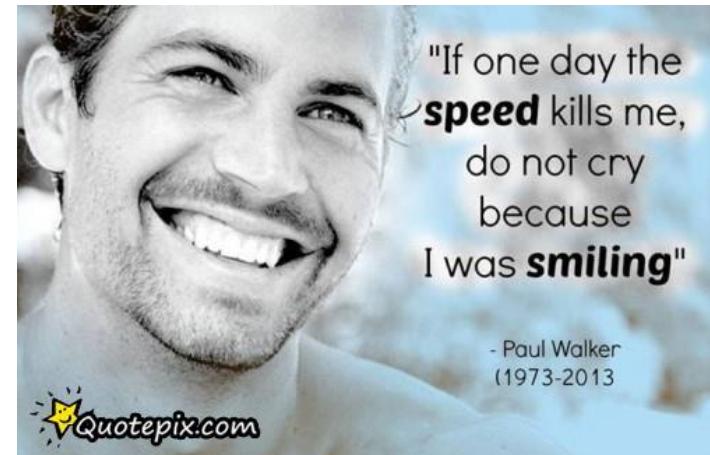
☛ Why? Pushing up the clock speed.

Why the clock needs to go faster

- Because it was there (?)
- A little googling tells you humans do obsess over speed
- Ignoring why, my job is to help, I'll talk about the what and how.

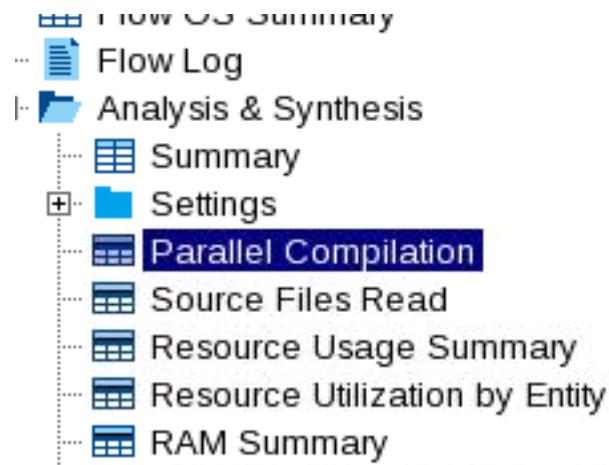


I FEEL THE NEED...
THE NEED
FOR SPEED.
-TOP GUN



There aren't a lot of quote images for “parallelism”

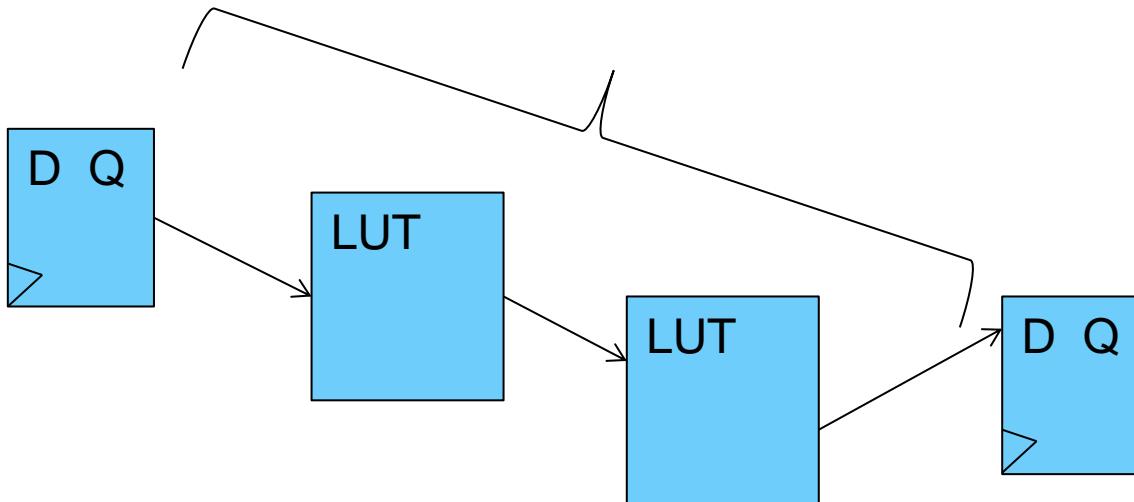
- Complicated, Amdahl.
- N things at 100 MHz is generally less exciting than one thing at N-hundred MHz.
- People who say otherwise can't do N-hundred (yet)



b	Usage by Processor	% Time Used
7	Processor 1	100.0%
1	Processors 2-7	4.3%
2	Processor 8	4.2%
3	Processors 9-10	4.0%
4	Processor 11	4.0%
5	Processors 12-13	3.9%
6	Processors 14-15	3.8%
7	Processor 16	3.8%
8		

Our pipeline context

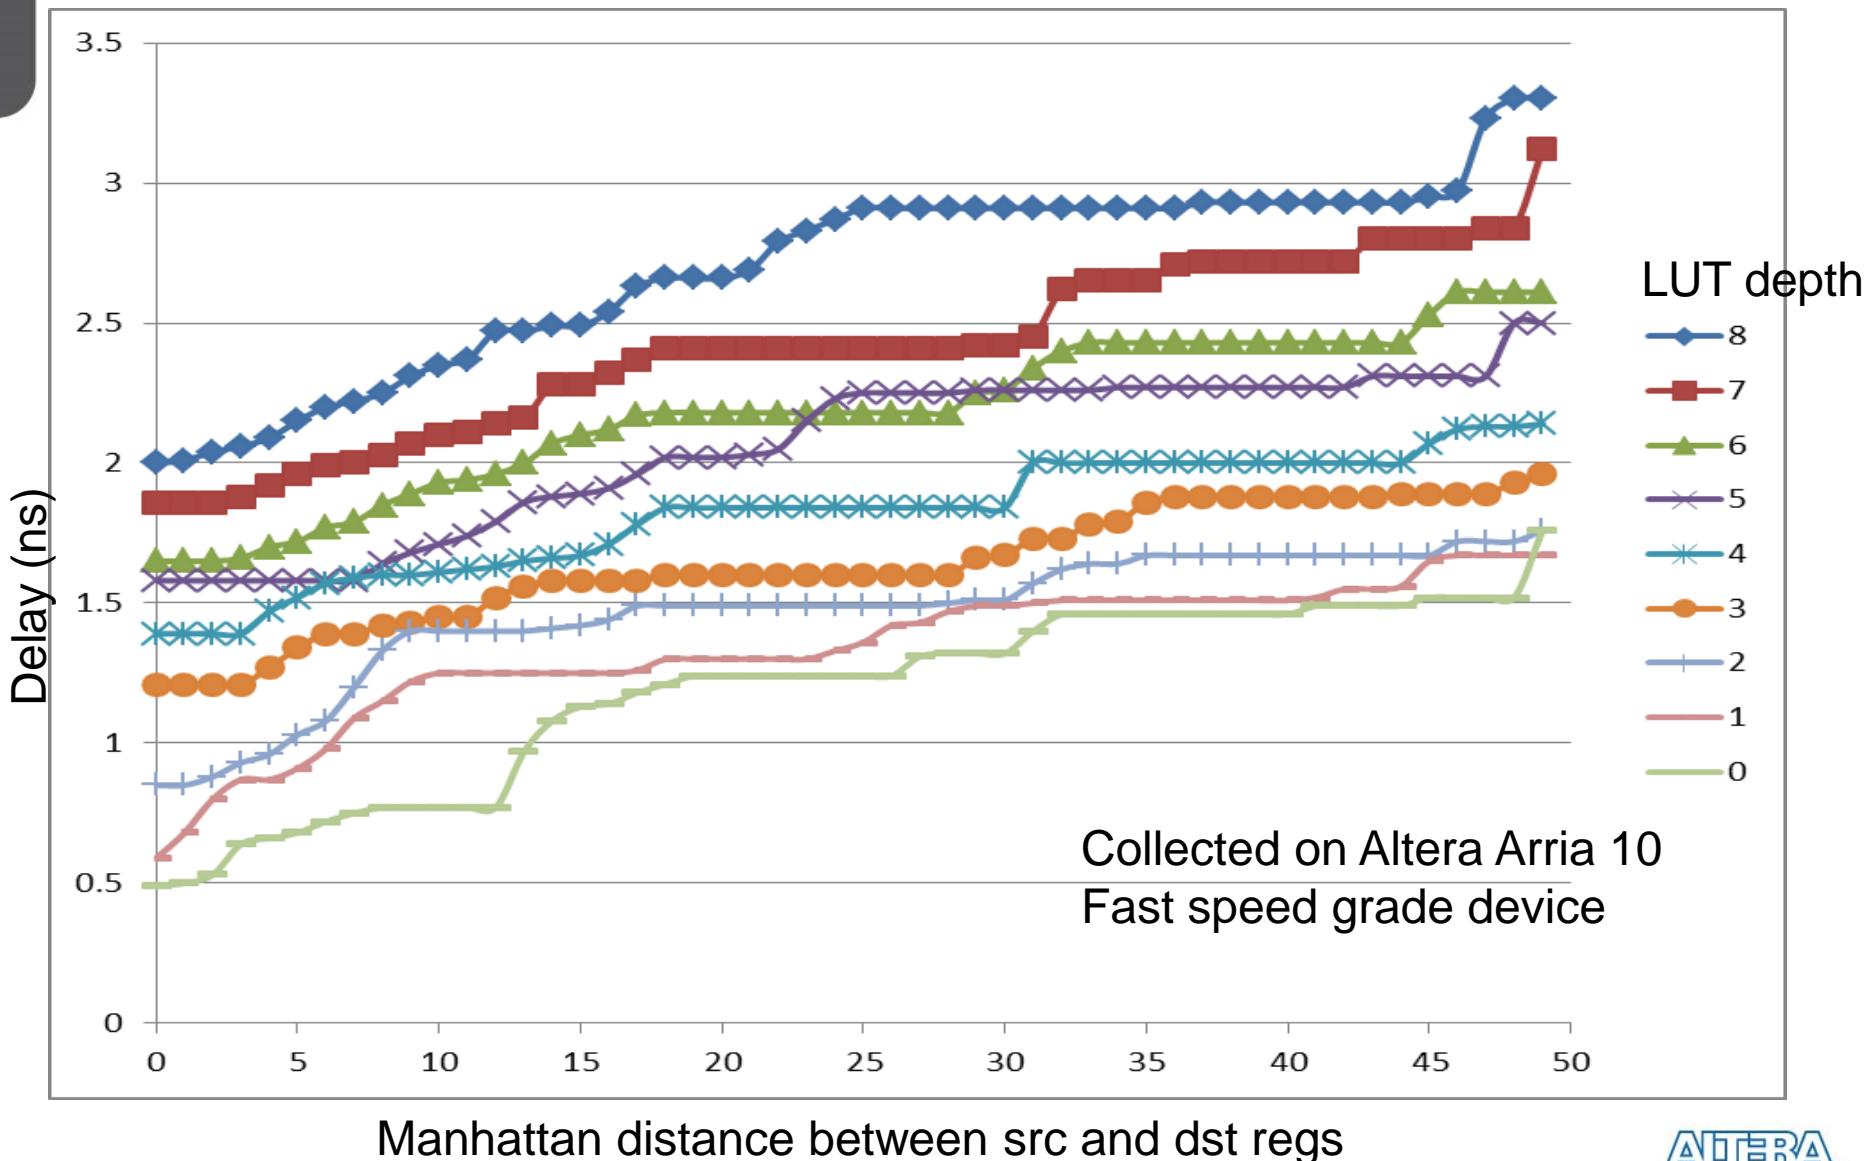
- In (mainstream) FPGAs the clock speed is governed by



The worst case of - How much look up table (LUT) and wire is between some pair of registers in the design

Myriad secondary effects

Delay for single path register -> (N) LUTS -> register



Interpreting these path delays

Respectable summary formula for Arria 10

- Delay = (188ps * LUT_depth) + (22ps * distance) + 540ps
- Pretty well aligned with silicon model delays

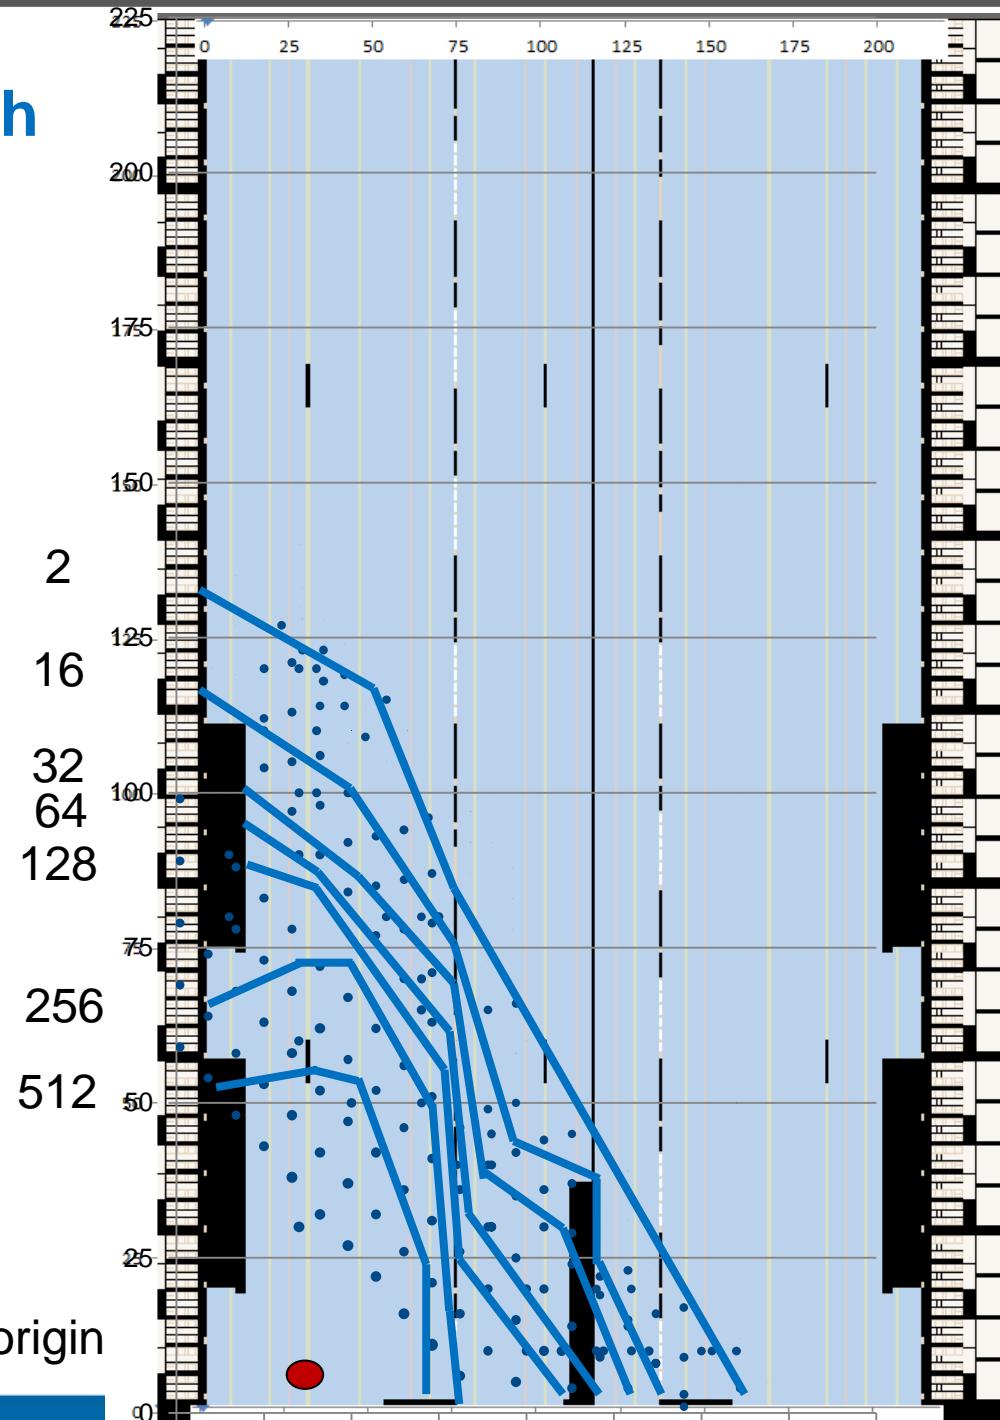
Can I really run no-LUT no-routing R2R at 1.8GHz ?

- Maybe. You can timing close a little FIFO at 800MHz
- On the desk, room temperature you can spin it up well past 1GHz before it fails.

Routing is dangerous. Traveling 8.5 points ~ = Another LUT delay

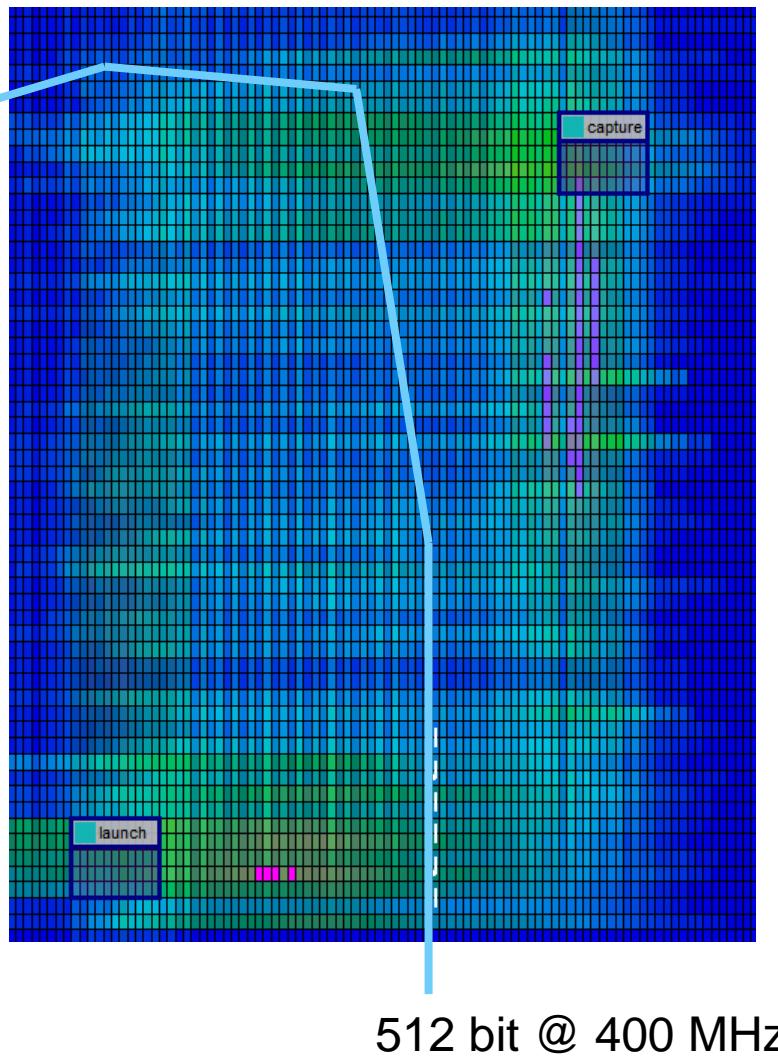
These are singleton near-optimal hops, let's look at repeating them, with zero LUTs

Bus Reach



How far can I travel between two N bit registers at 400MHz?

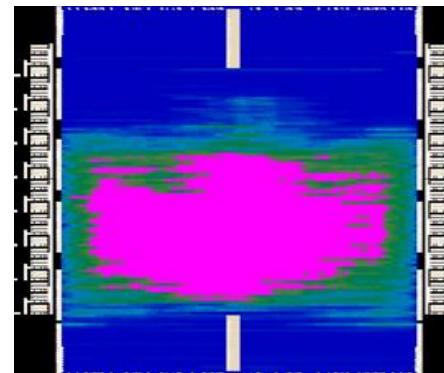
Bus reach sample point



- This is a zoomed in routing view of a missed 512 bit probe, it is only rated for 309 MHz
- You can see wires traversing all across the rectangle, and some emerging congestion in pink.
- If it were 16 bits it would be meeting 400MHz comfortably

What's going on

- ⌚ Competition for physical resources
- ⌚ Minimum of N trials
- ⌚ Widest busses in somewhat routine use today are 1024-1280 bit
 - This is 300-500G
 - Yeah it hurts ->
- ⌚ Around 2048 it becomes impractical to move



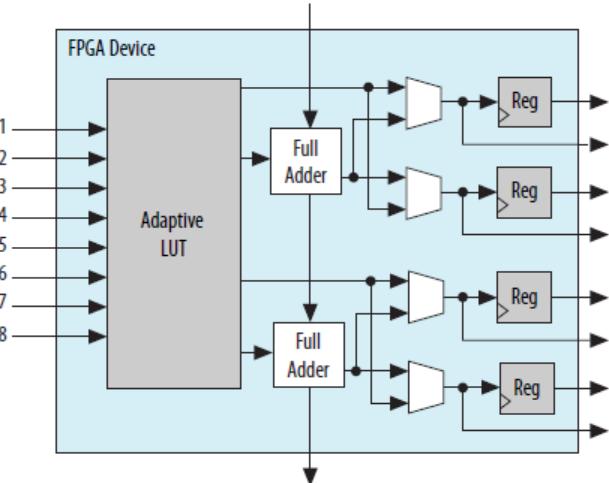
(This is a 400GE/500G link bridge on Stratix 5)

Key observations

- The LUTs are fundamentally fast. If you aren't travelling you can go through a bunch of them (e.g. a carry chain)
- The speed of travel is impeded by the number of travelers
- At typical communication distances (say 15) the routing delay and delay of a couple LUTs are about equal
- There is an unpleasant noise-floor term, traveling zero distance through zero LUTs costs a half nanosecond
 - Stratix 10 is going to tackle this aspect

Back to “hyper pipeline”

- ↳ The physical resources in the A10 device provide >1 register per LUT
 - ↳ Operate comfortably > 500 MHz
 - ↳ Deviating far from here is inefficient
-
- ↳ Taken together with the previous speed numbers you see a sweet spot where a 1:1 REG:LUT path can traverse about 50 points (1/4 device), operate at 500MHz, handle bus up to 512 bits
-
- ↳ This is more heavily pipelined than “average” FPGAs today, but routine for high speed interface logic



Define “high speed interface logic”

- ⬚ The label makes me uncomfortable
 - An AND gate in a state machine, or an ALU, or masking out a packet header at 400G. He's just an AND gate.
 - The ‘type’ of logic is I think a completely human red herring construct
- ⬚ Anyway, “the logic near the edges of the chip, usually running a well defined protocol like Ethernet, Interlaken, PCIE, memory control, usually talking to neighboring chips” is “interface logic”
- ⬚ It tends to be made by obsessive designers for reuse by others. As such it tends to be fast, compact, high bandwidth
- ⬚ Using Ethernet as an example...

If you want to understand 40/100/400 G Ethernet

- ☛ There is a 450 page IEEE specification which talks about the relevant portions of the standard Why so long?
Endless detail
- ☛ The Altera 100G Ethernet with all optional features is about 1000 pages of Verilog Parallelism, cell optimizations
- ☛ You can explain the digital meat pretty well with a 2 page C program, which talks mostly about byte operations, some 64 and 66 bit.

Bandwidth

<IDLE><SOP>Mary had a little lam

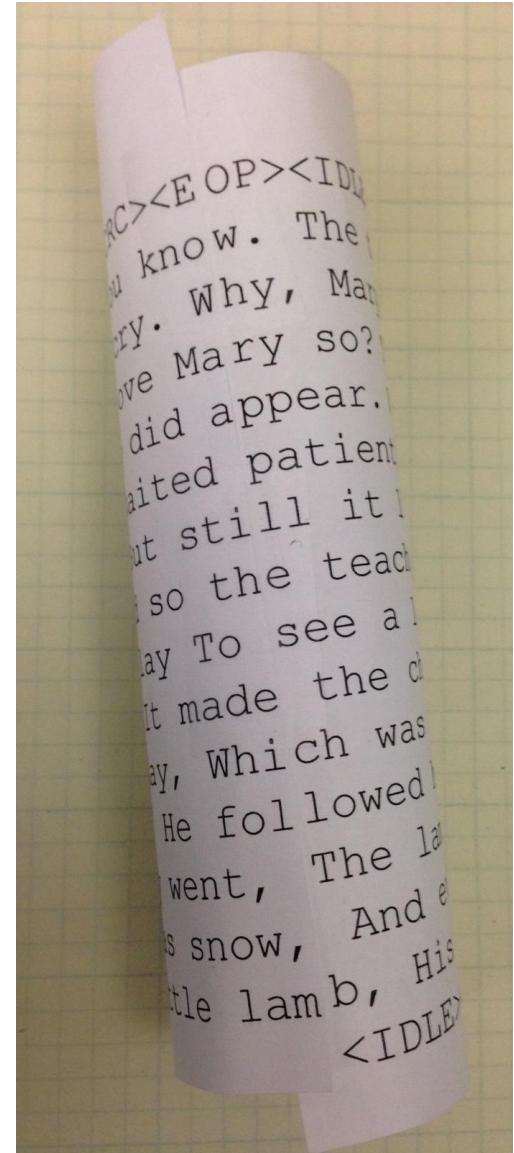
OP><IDLE>
w. The teacher did reply.<CRC><E
hy, Mary loves the lamb, you kno
ry so? The eager children cry. W
pear. Why does the lamb love Ma
patiently about, Till Mary did a
ll it lingered near, And waited
e teacher turned it out, But sti
see a lamb at school. And so th
the children laugh and play To
ch was against the rule, It made
lowed her to school one day, Whi
The lamb was sure to go. He fol
And everywhere that Mary went,
b, His fleece was white as snow,
<IDLE><SOP>Mary had a little lam

Concept – 8 bits at
12.5 GHz

Status Quo – 256 bits at
390.625 MHz

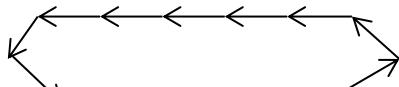
Why that hurts

- Ethernet is not one of the “Embarrassingly Parallel” problems, the circuit growth is typically violent
- Expansion currently done by humans (unfortunately) outperforming the HLS CAD tools
- The topology of the circuitry needs to follow the topology of the data, which doesn’t honor the break we introduced between “lamb” and “b”
 - It is not two dimensional =>
 - Imagine place and route tool’s problem view of this tube on saran wrap



Placement in plastic

OP><IDLE>
w. The teacher d
hy, Mary loves t
ry so? The eager
ppear. Why does
patiently about,
it lingered n
e teacher turned
see a lamb at sc
the children fa
ch was against t
fowed her to sch
The Lamb was sur
And everywhere
b, His fleece wa
<IDLE><SOP>Mary



Pretty good

OP><IDLE>

w. The teacher did reply.<CRC><E
hy, Mary loves the lamb, you kno
ry so? The eager children cry. W
ppear. Why does the lamb love Ma
patiently about, Till Mary did a
ll it lingered near, And waited
e teacher turned it out, But sti
see a lamb at school. And so th
the children laugh and play To
ch was against the rule, It made
lowed her to school one day, Whi
The lamb was sure to go. He fol
And everywhere that Mary went,
b, His fleece was white as snow,
<IDLE><SOP>Mary had a little lam



thrilled

FAQ

Does the place and route see that tube pattern?

- Often.
- Examine the directional routing in the chip planner, which is turbulent in areas where it should conceptually be one-way

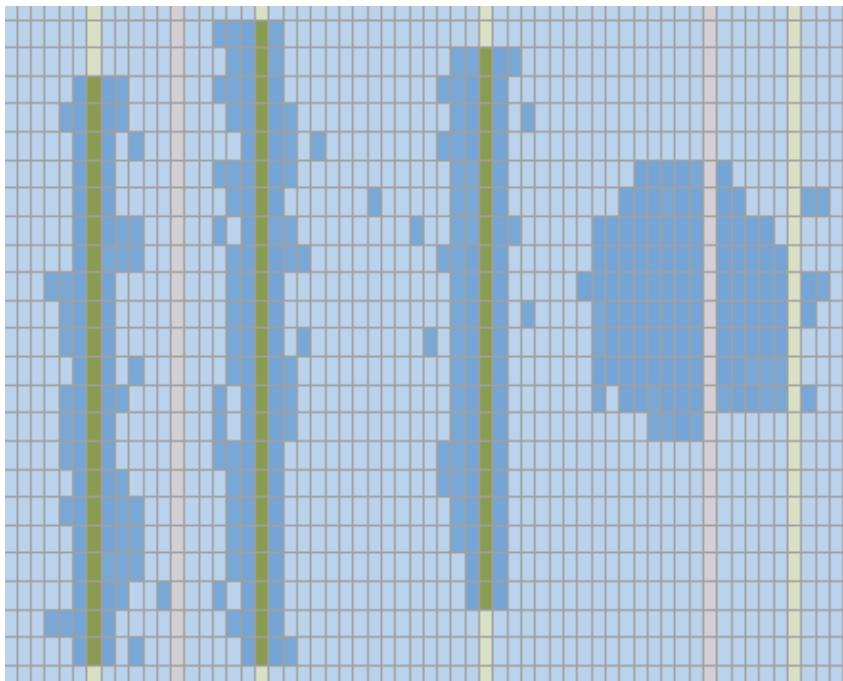
What happens when tubes intersect (e.g. a switch) ?

- Topological havoc.

What can the designer do to help keep speed up?

- Add more registers

Designer awareness of spread and grouping



- There are a variety of situations that force bits physically apart.
E.g. coupling to a RAM or SERDES pin
- There are a variety of computations that pull bits together in competing patterns.
E.g. bit 0 relates to other bits of a word, the previous cycle, and bit 0 of other words.
- If you scrutinize high speed designs you'll encounter extra pipeline stages and manual register duplication aimed at releasing this sort of tension.

Status quo of FPGA pipeline

- ⬚ There is a wide variety of FPGA circuitry in the world running heavily pipelined (1 or 2 LUTs deep) logic around the 256-512 bit bus, 300-400MHz range
 - Or 100-200 gigabit per second flows
- ⬚ And of course plenty of boutique logic above, legacy below, and variations
- ⬚ The speed drifts up slowly as the devices improve
 - But perpetually less than the appetite
- ⬚ There is always friction

#1 request

- ↳ “I am having trouble closing timing, please slow down by a factor of 2 and increase the bus width by a factor of 2”
- ↳ Recipe for disappointment
- ↳ Area and latency $> 2x$ (area $\sim=$ cost, compile time)
 - Poor inherent parallelism
 - Unpleasant surprises, like encountering multiple packets per cycle
- ↳ Flat power if lucky
- ↳ A lot of the new found cycle time evaporates into routing, double area implies double the *required* travel distance
- ↳ We help customers who ask for this, many revert back

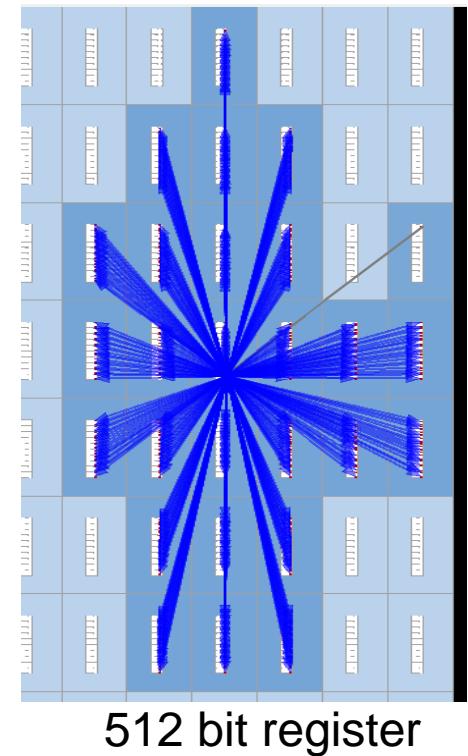
#2 request

“I want to go faster, but inserting registers isn’t helping anymore.”

Why?

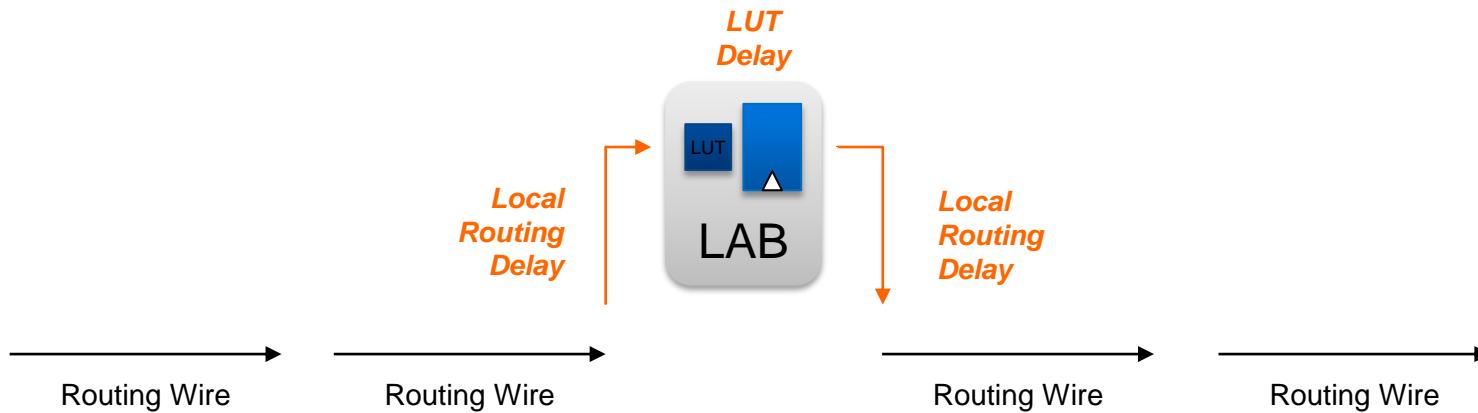
- A bus register takes up physical space which pushes things apart, which tends to make them slower
- It takes time to get in and out of a register recall the zero cost point to point path at a half nanosecond, not 0 nanoseconds

** Stratix 10 HyperFlex registers **



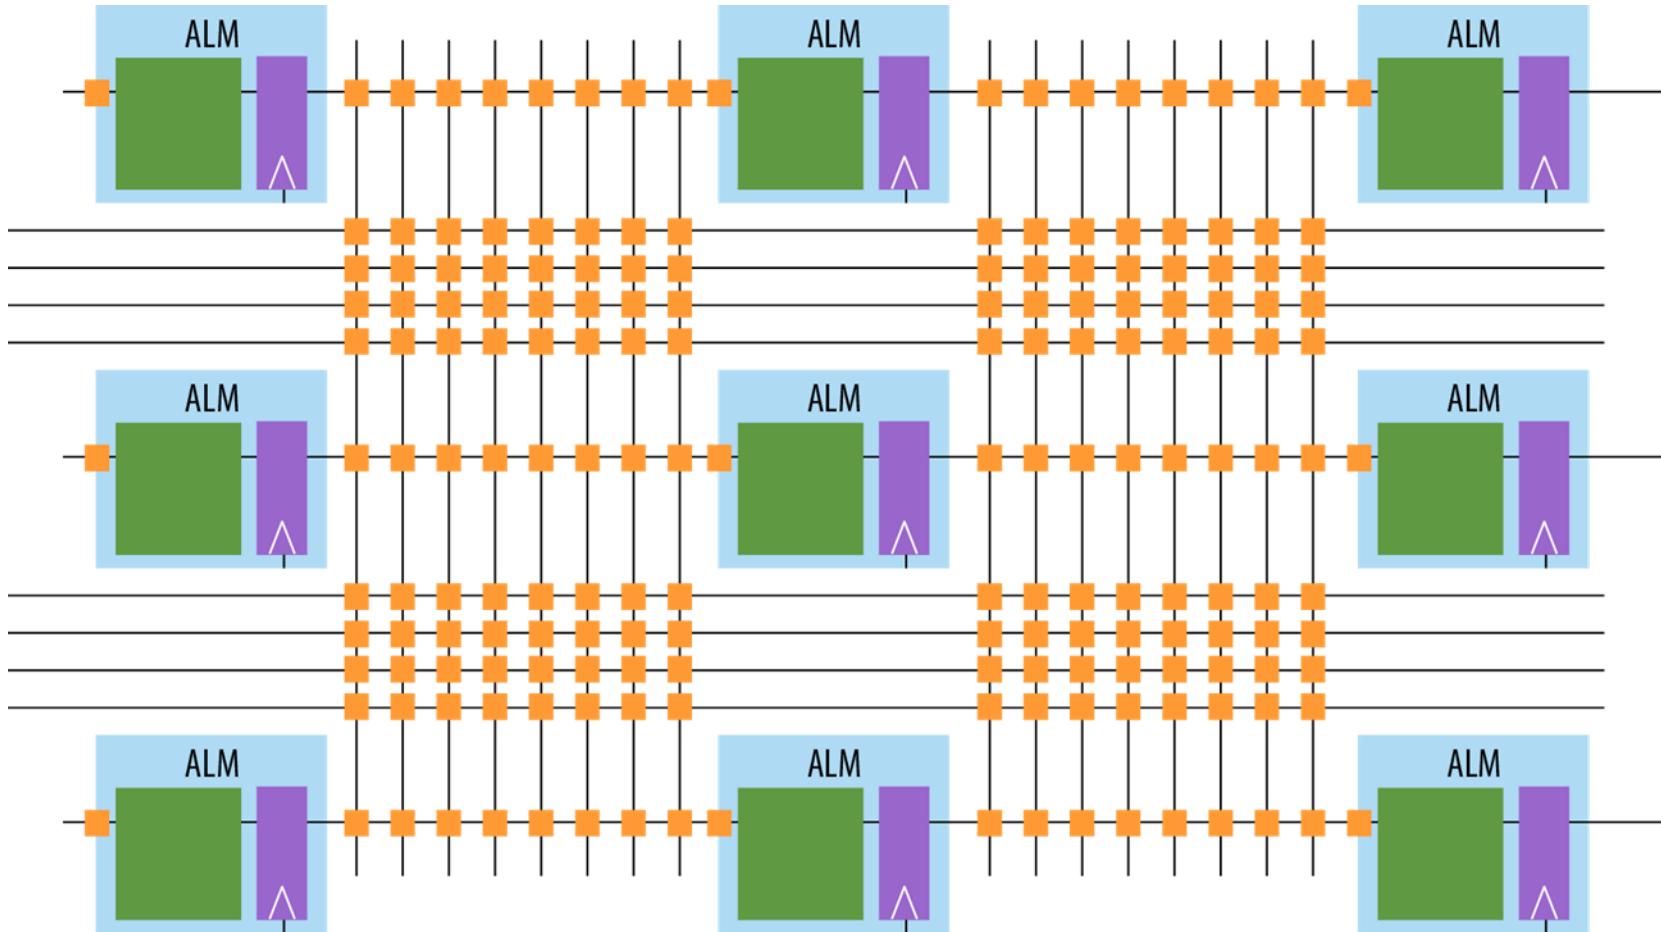
Why adding registers fizzles out at some point

- Fast designs today have more routing than logic delay
- It's annoying to get off the wire, find a home for a register, mux in and out again.

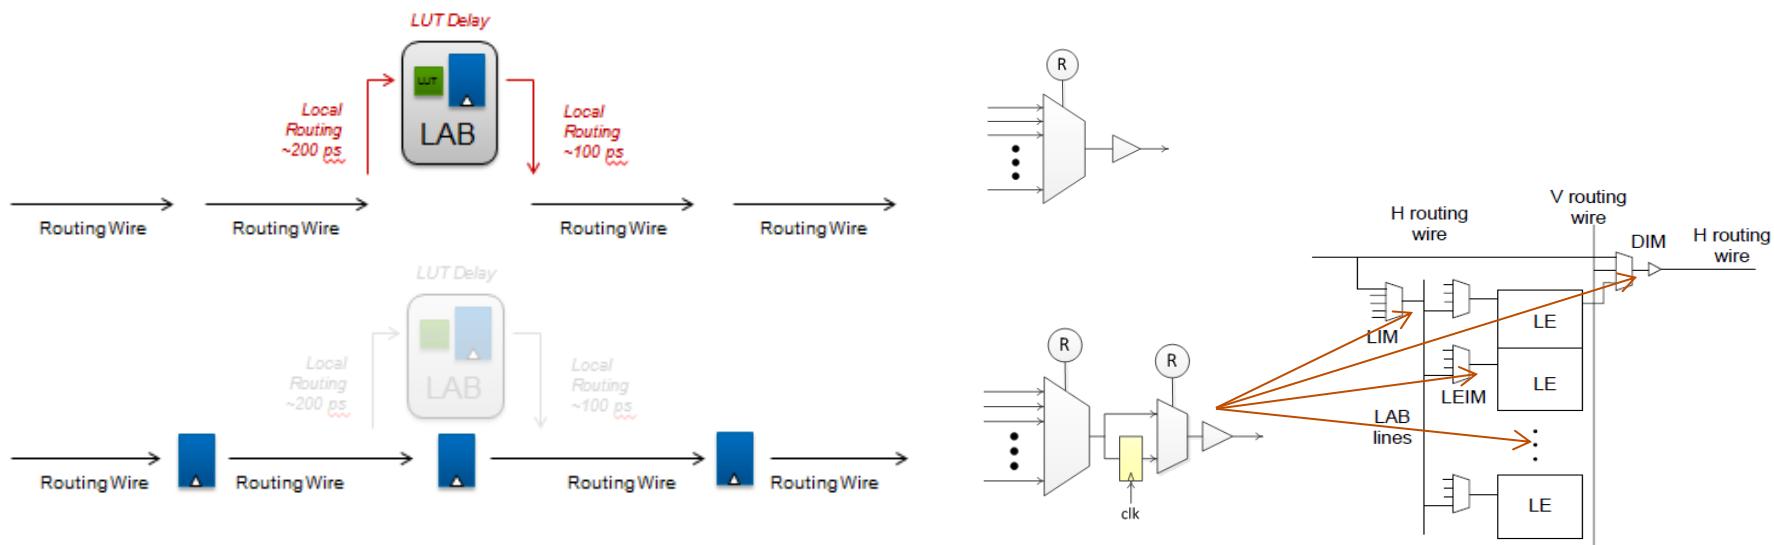


- Observation – the registers are TINY, far from datasheet scale

Stratix 10 has ‘plenty’ of Hyper-Registers distributed into the routing network



HyperFlex: Pipeline Registers by Design



- Routing muxes (all H/V wires) have *optional* registers
 - Including LAB, M20K and DSP block inputs, CC, SCLR/CE

Architectural Goals:

- Perfect balance – P&R chooses the right register (of many) to turn on
- Simple Software – Re-timing is a simple push/pull along the path
- No wasted LEs – Designs with high FF:LUT ratios no longer an issue
- No wasted routing – Don't have to route to find an available FF

What more registers do to a real circuit

Interlaken transmit datapath, 12 words (768 bit). From a Altera 300G hardware demo



itx_dp_12

Designed for 390 MHz plus 20% margin = 468 MHz on Stratix 5
Currently scoring 561 MHz on Arria 10



S10 target



Stratix 10 prototype
making 842 MHz

Variations assuming changes to register counts, hold avoidance software, etc..

The software (illegally) increased the latency from 13 to 14, and (legally) did some minor register retiming. The latency edit amounts to adding a layer of input registers and fixing up the valid schedule. Nothing extreme. Heavily pipelined logic doesn't need much rework to shine in the new fabric.

Implications of this result

- ☛ For a 10 minute RTL edit to legalize what the prototype software just assumed I would get (real) 842 MHz
- ☛ I listed the target at 781 MHz, which is double the current operating rate of 390 (390.625 MHz * 768 bit is 300Gbps)
- ☛ This thing is now a 600G Interlaken.
- ☛ If you actually wanted 300G
 - The datapath slices to 384 bit, area and latency (in ns) less than half
 - Stress taken off the place and route tool from half sized problem, speed will increase.
 - Option to move to a smaller device, option to add another port (speed is very convertible to other benefits)

Good secondary effects

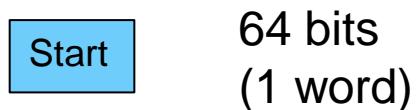
- Less parallel = easier to think about = less circuit, less buggy



1024 bits
(16 words)



512 bits
(8 words)



64 bits
(1 word)

- IP makers in the process of retargeting major cores
- Register retiming works dramatically better due to having more and finer grained choices

Challenges

- ☛ Harder to shake off small absolute delay surprises
- ☛ Unsettling for many FPGA designers to see 10x registers and be asked to not fill them up.
- ☛ Chip is similar scale, and capable of sustaining much higher toggle rates. Customer bandwidth appetite ~unlimited. Clock and power distribution more exciting.
- ☛ The usual concerns about verification and debug-ability in the presence of retiming

Where we're heading with soft IP

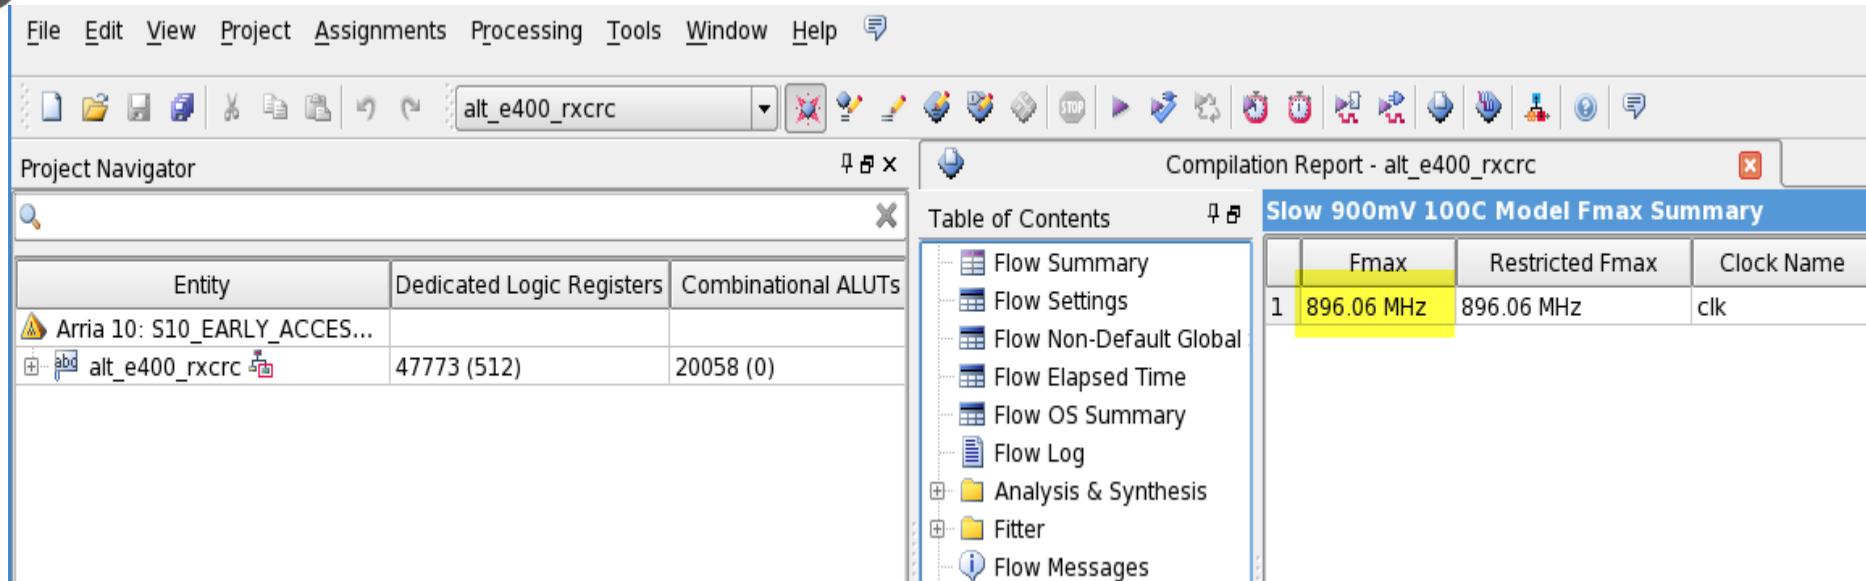
- Science projects at extreme speed, this is a Ethernet MAC targeted at 1.25GHz, described at Ethernet Technology Summit

Project	Device	Logic Cells (ALMs)	Registers	Performance (MHz)	Latency (ns)
10G Ethernet MAC	Arria 10	1500	2600	340	59
Falcon 10G MAC	Stratix 10	338	1231	1358	12.8

- Double pumping RAM / DSP / switching variants

Where we're heading with soft IP

- More ports and fatter pipes.



- This is a 400G Ethernet CRC assembly, used in Stratix 5 hardware demo on YouTube (390 MHz x 1024 bit)
- With minimal labor, it is modeling at 896 MHz
- 896 MHz * 1024 bit is a stone's throw from Terabit.

Where we're heading with soft IP

- ↳ In the presence of a register surplus, and fluid retiming, loops inevitably become critical.
- ↳ We're having a resurgence of interest in Shannon-izing decisions, methods for restructuring accumulations (both manually and automatically) and adapting early decisions to anticipated future retiming.

In summary

- ☛ Bandwidth is going to keep climbing, pipelines will grow steadily deeper, and architectural changes will take the sting out of the increased register count
- ☛ For more information
 - On www.altera.com – products – Stratix 10
 - ☛ Myriad white papers and app notes.
 - The training section is a (somewhat hidden) gem
 - ☛ <https://www.altera.com/products/fpga/stratix-series/stratix-10/support.html#training>

Lots of direct and detailed discussion written by SW / TS / IP staff on the RTL strategies they use to make FPGA designs go faster, with lab examples

Thank You